

(<https://databricks.com>)

# MEMUDU Alimatou Sadia

## Big Data Project

### Dataset Description

**Dataset:** <https://www.kaggle.com/datasets/shelvigarg/credit-card-buyers>

This dataset is a zip file composed of 2 csv files, test data credit card.csv and train data credit card.csv however only the train data credit card.csv will be used for this project.

Happy Customer Bank is a mid-sized private bank that deals in all kinds of banking products, like Savings accounts, Current accounts, investment products, credit products, among other offerings. The bank also cross-sells products to its existing customers and to do so they use different kinds of communication like tele-calling, e-mails, recommendations on net banking, mobile banking, etc. In this case, the Happy Customer Bank wants to cross sell its credit cards to its existing customers. The bank has identified a set of customers that are eligible for taking these credit cards. Now, the bank is looking for your help in identifying customers that could show higher intent towards a recommended credit card.

### Data Variables

ID: Unique Identifier for a row

Gender: Gender of the Customer

Age: Age of the Customer (in Years)

Region\_Code: Code of the Region for the customers

Occupation: Occupation Type for the customer

Channel\_Code: Acquisition Channel Code for the Customer (Encoded)

Vintage: Vintage for the Customer (In Months)

Credit\_Product: If the Customer has any active credit product (Home loan, Personal loan, Credit Card etc.)

AvgAccountBalance: Average Account Balance for the Customer in last 12 Months

Is\_Active: If the Customer is Active in last 3 Months

Is\_Lead(Target): If the Customer is interested for the Credit Card

0 : Customer is not interested

1 : Customer is interested

## Goal:

Build a full machine learning pipeline to predict if the customers will be interested in getting a credit card using PySpark.

# Data Collection

```
# Importing libraries
```

```
import pyspark
from pyspark.sql import SparkSession
from pyspark import pandas as pd
from pyspark.sql.functions import col, isnan, when, count
from pyspark.sql.functions import count_distinct
from pyspark.sql import SparkSession
from pyspark import pandas as pd
from pyspark.ml import Pipeline
from pyspark.ml.feature import
(OneHotEncoder, StandardScaler, StringIndexer, VectorAssembler)
from pyspark.ml.stat import Correlation
from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit
from pyspark.ml.classification import LogisticRegression, GBClassifier,
RandomForestClassifier
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.tuning import CrossValidator, CrossValidatorModel,
ParamGridBuilder
```

```
#pip install kaggle
#import kaggle
```

```
# copying my kaggle.json file to kaggle local folder
# src_path = r"C:\Users\Alimat sadia\Downloads\kaggle.json"
# dst_path = r"C:\Users\Alimat sadia\.kaggle\kaggle.json"
# shutil.copy(src_path, dst_path)
# print('Copied')

# dbutils.fs.cp("dbfs:///FileStore/tables/kaggle_token/kaggle.json",
"file:/root/.kaggle/")
# !kaggle datasets download -d shelvigarg/credit-card-buyers
# # # unzipping data

# !unzip /databricks/driver/credit-card-buyers.zip
```

```
spark=SparkSession.builder.getOrCreate()
```

```
spark
```

**SparkSession - hive**

**SparkContext**

[Spark UI](#)

Version

v3.3.0

Master

local[8]

AppName

Databricks Shell

# Data Exploration

```
# Read data file
```

```
df = spark.read.csv("/FileStore/tables/dataset/train_data_credit_card.csv",  
header=True, inferSchema=True)
```

```
# to check all the columns and their datatype
```

```
df.printSchema()
```

```
root
```

```
|-- ID: string (nullable = true)  
|-- Gender: string (nullable = true)  
|-- Age: integer (nullable = true)  
|-- Region_Code: string (nullable = true)  
|-- Occupation: string (nullable = true)  
|-- Channel_Code: string (nullable = true)  
|-- Vintage: integer (nullable = true)  
|-- Credit_Product: string (nullable = true)  
|-- Avg_Account_Balance: integer (nullable = true)  
|-- Is_Active: string (nullable = true)  
|-- Is_Lead: integer (nullable = true)
```

```
#Number of rows and columns
```

```
print("The shape of the data is ",(df.count(), len(df.columns)))
```

```
The shape of the data is (245725, 11)
```

```
# statistical summary of the data columns
```

```
df.describe().toPandas().transpose()
```

	0	1	2	3
<b>summary</b>	count	mean	stddev	min
<b>ID</b>	245725	None	None	222A8XWS
<b>Gender</b>	245725	None	None	Female
<b>Age</b>	245725	43.85630684708516	14.828671804648	23
<b>Region_Code</b>	245725	None	None	RG250
<b>Occupation</b>	245725	None	None	Entrepreneur
<b>Channel_Code</b>	245725	None	None	X1
<b>Vintage</b>	245725	46.95914131651236	32.35313570875409	7
<b>Credit_Product</b>	216400	None	None	No
<b>Avg_Account_Balance</b>	245725	1128403.1010194323	852936.3560692756	20790
<b>Is_Active</b>	245725	None	None	No
<b>Is_Lead</b>	245725	0.23720826126767727	0.4253718824871881	0

```
# viewing dataset
df.show(15)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|      ID|Gender|Age|Region_Code|  Occupation|Channel_Code|Vintage|Credit_P
roduct|Avg_Account_Balance|Is_Active|Is_Lead|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|NNVBBKZB|Female| 73|      RG268|      Other|      X3|      43|
No|      1045696|      No|      0|
|IDD62UNG|Female| 30|      RG277|    Salaried|      X1|      32|
No|      581988|      No|      0|
|HD3DSEMC|Female| 56|      RG268|Self_Employed|      X3|      26|
No|      1484315|      Yes|      0|
|BF3NC7KV|  Male| 34|      RG270|    Salaried|      X1|      19|
No|      470454|      No|      0|
|TEASRWXV|Female| 30|      RG282|    Salaried|      X1|      33|
No|      886787|      No|      0|
|ACUTYTWS|  Male| 56|      RG261|Self_Employed|      X1|      32|
No|      544163|      Yes|      0|
|ETQCZFEJ|  Male| 62|      RG282|      Other|      X3|      20|
null|      1056750|      Yes|      1|
|JJNJUQMQ|Female| 48|      RG265|Self_Employed|      X3|      13|
```

```
# Displaying uniques values of each column
```

```
for col_name in df.columns:
    unique_val = df.select(col_name).distinct().collect()
    print(f" {col_name}")
    print(f"Unique values count: {len(unique_val)}")
```

```
ID
Unique values count: 245725
Gender
Unique values count: 2
Age
Unique values count: 63
Region_Code
Unique values count: 35
Occupation
Unique values count: 4
Channel_Code
Unique values count: 4
Vintage
Unique values count: 66
Credit_Product
Unique values count: 3
Avg_Account_Balance
Unique values count: 135292
Is_Active
Unique values count: 2
Is_Lead
```

## Number of unique value per columns

```
# checking for missing value
```

```
# Find Count of Null, None, NaN of All DataFrame Columns
df.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c in
df.columns]
).show()
```

```
+---+-----+---+-----+-----+-----+-----+-----+-----+
| ID|Gender|Age|Region_Code|Occupation|Channel_Code|Vintage|Credit_Product|Avg_Account_Balance|Is_Active|Is_Lead|
+---+-----+---+-----+-----+-----+-----+-----+-----+
|  0|      0|  0|          0|          0|          0|          0|          0|          29325|          0|          0|
+---+-----+---+-----+-----+-----+-----+-----+-----+
|  0|      0|      0|          0|          0|          0|          0|          0|          29325|          0|          0|
+---+-----+---+-----+-----+-----+-----+-----+-----+
```

```
-----+-----+-----+
```

Credit Product has the highest number of missing value and is the only column with missing value

```
# Value counts of Credit_Products columns
```

```
df.groupBy('Credit_Product').count().orderBy('count',
ascending=False).show()
```

```
+-----+-----+
|Credit_Product| count|
+-----+-----+
|              No|144357|
|              Yes| 72043|
|              null| 29325|
+-----+-----+
```

Now we can see the most frequent value of this column is No, so we can replace the missing values by the most frequent one instead of deleting the rows

```
df = df.fillna({"Credit_Product": 'No'})
df.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|      ID|Gender|Age|Region_Code|  Occupation|Channel_Code|Vintage|Credit_P
roduct|Avg_Account_Balance|Is_Active|Is_Lead|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|NNVBBKZB|Female| 73|      RG268|      Other|      X3|      43|
No|      1045696|      No|      0|
|IDD62UNG|Female| 30|      RG277|      Salaried|      X1|      32|
No|      581988|      No|      0|
|HD3DSEMC|Female| 56|      RG268|Self_Employed|      X3|      26|
No|      1484315|      Yes|      0|
|BF3NC7KV|  Male| 34|      RG270|      Salaried|      X1|      19|
No|      470454|      No|      0|
|TEASRWXV|Female| 30|      RG282|      Salaried|      X1|      33|
No|      886787|      No|      0|
|ACUTYTWS|  Male| 56|      RG261|Self_Employed|      X1|      32|
No|      544163|      Yes|      0|
|ETQCZFEJ|  Male| 62|      RG282|      Other|      X3|      20|
No|      1056750|      Yes|      1|
|JJNJUQMQ|Female| 48|      RG265|Self_Employed|      X3|      13|
```

```
df.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c in
df.columns]
).show()
```

```
+---+-----+---+-----+-----+-----+-----+-----+
-----+-----+-----+
| ID|Gender|Age|Region_Code|Occupation|Channel_Code|Vintage|Credit_Product|A
vg_Account_Balance|Is_Active|Is_Lead|
+---+-----+---+-----+-----+-----+-----+-----+
-----+-----+-----+
|  0|      0|  0|          0|          0|          0|      0|          0|
0|      0|      0|
+---+-----+---+-----+-----+-----+-----+-----+
-----+-----+-----+
```

Great News ! our data is free of missing values

## Information Extraction With Spark SQL

# we can easily register a DataFrame as a table

```
df.createOrReplaceTempView("data")
```

# Distinct region

```
df.select(
    count_distinct("Region_Code").alias("distinct region code")
).show()
```

```
+-----+
|distinct region code|
+-----+
|                  35|
+-----+
```

```
spark.sql("SELECT * FROM data").show()
```

```
+-----+-----+---+-----+-----+-----+-----+-----+
-----+-----+-----+
|      ID|Gender|Age|Region_Code|  Occupation|Channel_Code|Vintage|Credit_P
roduct|Avg_Account_Balance|Is_Active|Is_Lead|
```



[illegible]

# Number of male with active credit card

```
spark.sql("SELECT count( DISTINCT ID) as Number_of_male_active FROM data
where data.Credit_Product='Yes' and data.Gender='Male']").show()
```

Number_of_male_active
41738

# Percentage of male and female users who has been Active in the last 3 Months

```
spark.sql("SELECT data.Gender,ROUND((SUM(CASE WHEN data.Is_Active='Yes' THEN 1 ELSE 0 END) / COUNT(*) ) * 100,2) as percentage FROM data GROUP BY data.Gender").show()
```

Gender	percentage
Female	35.41
Male	41.69

```
# Target variable Is_Lead Distribution
```

```
spark.sql("SELECT data.Is_Lead, COUNT(data.Is_lead) as distinct ,  
(COUNT(data.Is_lead) / (SELECT COUNT(data.Is_lead) FROM data))*100 as  
percentage FROM data GROUP BY data.Is_Lead").show()
```

```

+-----+-----+-----+
|Is_Lead|distinct|percentage|
+-----+-----+-----+
|      1|    58288|23.720826126767726|
|      0|   187437|76.27917387323228|
+-----+-----+-----+

```

Our data is imbalanced, the number of 1 is more larger than the number of 0.

```

# Average account balance per region
spark.sql("SELECT data.Region_Code , (AVG(data.Avg_Account_Balance)) as
Average_account_balance from data Group By data.Region_Code order by
Average_account_balance desc").show()

```

```

+-----+-----+-----+
|Region_Code|Average_account_balance|
+-----+-----+-----+
|      RG283|1475574.2942276313|
|      RG284|1473165.5680641823|
|      RG268|1463900.0817610063|
|      RG254|1407392.135394933|
|      RG253|1374248.0543595264|
|      RG262|1200158.0458612975|
|      RG276|1059816.6505065123|
|      RG269|1027003.1523591505|
|      RG274|996307.0546727204|
|      RG277|983289.3474972711|
|      RG261|981286.4003668282|
|      RG282|953867.7062961056|
|      RG278|913786.0779363337|
|      RG281|889643.6980168859|
|      RG272|885223.5024752475|
|      RG255|868249.742814668|
|      RG257|858105.4543517457|
|      RG273|856303.2921947965|

```

RG283 is the region with highest average account balance

## Data Preprocessing

The ID column is not necessary in this data let's remove it

```
df= df.drop("ID")
df.show()
```

```
+-----+---+-----+-----+-----+-----+-----+
|Gender|Age|Region_Code|Occupation|Channel_Code|Vintage|Credit_Product|Av
g_Account_Balance|Is_Active|Is_Lead|
+-----+---+-----+-----+-----+-----+-----+
|Female| 73|      RG268|      Other|      X3|    43|      No|
1045696|    No|      0|
|Female| 30|      RG277|    Salaried|      X1|    32|      No|
581988|    No|      0|
|Female| 56|      RG268|Self_Employed|      X3|    26|      No|
1484315|    Yes|      0|
|  Male| 34|      RG270|    Salaried|      X1|    19|      No|
470454|    No|      0|
|Female| 30|      RG282|    Salaried|      X1|    33|      No|
886787|    No|      0|
|  Male| 56|      RG261|Self_Employed|      X1|    32|      No|
544163|    Yes|      0|
|  Male| 62|      RG282|      Other|      X3|    20|      No|
1056750|    Yes|      1|
|Female| 48|      RG265|Self_Employed|      X3|    13|      No| //
```

```
#Float type columns of this data frame are:
```

```
target_col='Is_Lead'
```

```
numerical_cols = [df.dtypes[i][0] for i in range(len(df.dtypes)) if
df.dtypes[i][1]=='int']
```

```
numerical_cols.pop(-1)
```

```
numerical_cols
```

```
Out[23]: ['Age', 'Vintage', 'Avg_Account_Balance']
```

```
#Define string-type columns:
```

```
string_cols = [df.dtypes[i][0] for i in range(len(df.dtypes)) if
df.dtypes[i][1]=='string']
```

```
string_cols
```

```
Out[24]: ['Gender',
'Region_Code',
'Occupation',
'Channel_Code',
'Credit_Product',
'Is_Active']
```

In order to feed this data to a model we have to transform the categorical features into number using OneHotEncoder

```
# string columns encoding

string_cols_indexer = [f"{string_col}_indexer" for string_col in
string_cols]
string_cols_vector = [f"{string_col}_vector" for string_col in
string_cols_indexer]

string_cols_encode_stage= [
    StringIndexer(
        inputCols=string_cols,
        outputCols=string_cols_indexer,
    ),
    OneHotEncoder(
        inputCols=string_cols_indexer,
        outputCols=string_cols_vector)]

# numerical data scaling

numerical_scale_stage = [
    VectorAssembler(
        inputCols=numerical_cols,
        outputCol="assembled_numeric_cols",
    ),
    StandardScaler(
        inputCol="assembled_numeric_cols",
        outputCol="scaled_numeric_cols",
    ),
]

# group all the features into one columns

features_assembler = [
    VectorAssembler(
        inputCols=string_cols_vector + ["scaled_numeric_cols"],
        outputCol="features",
    )
]
```

```
target_stage = [StringIndexer(inputCol=target_col, outputCol="label")]
```

```
# Display effect on the whole data
```

```
Pipeline(stages=target_stage).fit(df).transform(df).show(5)
```

```
+-----+---+-----+-----+-----+-----+-----+-----+
|Gender|Age|Region_Code|Occupation|Channel_Code|Vintage|Credit_Product|Av
g_Account_Balance|Is_Active|Is_Lead|label|
+-----+---+-----+-----+-----+-----+-----+-----+
|Female| 73|      RG268|      Other|      X3|      43|      No|
1045696|      No|      0|      0.0|
|Female| 30|      RG277|      Salaried|      X1|      32|      No|
581988|      No|      0|      0.0|
|Female| 56|      RG268|Self_Employed|      X3|      26|      No|
1484315|      Yes|      0|      0.0|
|  Male| 34|      RG270|      Salaried|      X1|      19|      No|
470454|      No|      0|      0.0|
|Female| 30|      RG282|      Salaried|      X1|      33|      No|
886787|      No|      0|      0.0|
+-----+---+-----+-----+-----+-----+-----+-----+
+-----+---+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
# Splitting the data into train and test data
```

```
train, test = df.randomSplit(weights=[0.8, 0.2], seed=42)
print(f"Number of data in the train set: {train.count()}")
print(f"Number of data in the test set: {test.count()}")
```

```
Number of data in the train set: 196584
```

```
Number of data in the test set: 49141
```

```
# combining all the stages displayed above into a single pipeline
```

```
pipe = Pipeline(stages=string_cols_encode_stage + numerical_scale_stage +
features_assembler + target_stage)
data_preprocessing = pipe.fit(train)
```

```
data_preprocessing_train =
```

```
data_preprocessing.transform(train).select("features", "label")
```

```
data_preprocessing_train.show(5, truncate=False)
```

```

+-----+
| features |
| label |
+-----+
+-----+
| (46,[12,35,40,42,43,44,45],[1.0,1.0,1.0,1.0,1.5527608114781,0.4327398259957207,1.1804160076654502]) | 0.0 |
| (46,[26,36,38,41,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4018298384245978,0.4340912254307288]) | 0.0 |
| (46,[26,36,38,42,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4018298384245978,1.5836082643242448]) | 0.0 |
| (46,[26,36,38,41,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4327398259957207,0.6636636693035469]) | 0.0 |
| (46,[26,36,38,42,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4327398259957207,0.7558367647533992]) | 0.0 |
+-----+
+-----+
only showing top 5 rows

```

```

data_preprocessing_test =
data_preprocessing.transform(test).select("features", "label")
data_preprocessing_test.show(5, truncate=False)

```

```

+-----+
| features |
| label |
+-----+
+-----+
| (46,[26,36,38,42,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4018298384245978,0.6337383310425381]) | 0.0 |
| (46,[26,36,38,42,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4636498135668436,0.6812042865299271]) | 0.0 |
| (46,[26,36,38,42,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4636498135668436,0.9455009278527056]) | 0.0 |
| (46,[11,36,38,41,43,44,45],[1.0,1.0,1.0,1.0,1.6202721511075826,0.4327398259957207,0.6210376792389842]) | 0.0 |
| (46,[17,37,38,41,42,43,44,45],[1.0,1.0,1.0,1.0,1.0,1.6202721511075826,0.4327398259957207,1.2018123896293065]) | 0.0 |
+-----+
+-----+
only showing top 5 rows

```

# Model Building and Cross Validation

## Logistic Regression

```
lr_model = LogisticRegression()

# hyperparameter tuning
paramGrid = (
    ParamGridBuilder()
    .addGrid(lr_model.maxIter, [50, 100, 150])
    .addGrid(lr_model.regParam, [0.0, 0.5, 1.0])
    .addGrid(lr_model.elasticNetParam, [0.0, 0.5, 1.0])
    .build()
)

# Create a CrossValidator
evaluator = BinaryClassificationEvaluator()

cv = CrossValidator(
    estimator=lr_model,
    estimatorParamMaps=paramGrid,
    evaluator=evaluator,
    parallelism=8,
    numFolds=5,
)

cross_val_model = cv.fit(data_preprocessing_train)
```

```
# Extract the best model
best_model_lr = cross_val_model.bestModel
print(best_model_lr)

predictions_lr=best_model_lr.transform(data_preprocessing_test)
# evaluate best result result
print( "Logistic regression accuracy score on the test
set:",evaluator.evaluate(predictions_lr),"\\n")

print("Logistic Regression Prediction Table")
predictions_lr.select('label', 'prediction', 'probability').show(20)
```

LogisticRegressionModel: uid=LogisticRegression\_73cebbe1b8bc, numClasses=2, numFeatures=46  
Logistic regression accuracy score on the test set: 0.7341620461882972

Logistic Regression Prediction Table

label	prediction	probability
0.0	0.0	[0.89483941713461...
0.0	0.0	[0.89282481916175...
0.0	0.0	[0.89285797885386...
0.0	0.0	[0.90262157197535...
0.0	0.0	[0.96799815751839...
0.0	0.0	[0.91890817892784...
0.0	0.0	[0.91673685764783...
0.0	0.0	[0.96909688418369...
0.0	0.0	[0.89521304444288...
0.0	0.0	[0.89528125833428...
0.0	0.0	[0.92154765495802...
0.0	0.0	[0.91794971142531...
0.0	0.0	[0.91798555104292...

## Gradient-Boosted Tree

```
gbt_model = GBTClassifier()

# hyperparameter tuning
paramGrid = (ParamGridBuilder()
              .addGrid(gbt_model.maxDepth, [2, 4, 6])
              .addGrid(gbt_model.maxIter, [10, 15, 20])
              .build())

# Create a CrossValidator
evaluator = BinaryClassificationEvaluator()
```



```

cv = CrossValidator(
    estimator=gbt_model,
    estimatorParamMaps=paramGrid,
    evaluator=evaluator,
    parallelism=8,
    numFolds=5,
)

cross_val_model = cv.fit(data_preprocessing_train)

# Extract the best model
best_model_gbt = cross_val_model.bestModel
print(best_model_gbt)

predictions_gbt=best_model_gbt.transform(data_preprocessing_test)

# evaluate best result result
print( "GBTCClassifier accuracy score on the test
set:",evaluator.evaluate(predictions_gbt),"\n")

print("GBTCClassifier Prediction Table")
predictions_gbt.select('label', 'prediction', 'probability').show(20)

GBTCClassificationModel: uid = GBTCClassifier_df12f6811d27, numTrees=20, numCl
asses=2, numFeatures=46
GBTCClassifier accuracy score on the test set: 0.7888099034858672

GBTCClassifier Prediction Table
+-----+-----+-----+
|label|prediction|          probability|
+-----+-----+-----+
|  0.0|          0.0|[0.89088839340541...|
|  0.0|          0.0|[0.88507501011413...|
|  0.0|          0.0|[0.88698151346515...|
|  0.0|          0.0|[0.91464453799921...|
|  0.0|          0.0|[0.93863371810767...|
|  0.0|          0.0|[0.91464453799921...|
|  0.0|          0.0|[0.91464453799921...|
|  0.0|          0.0|[0.94204513281611...|
|  0.0|          0.0|[0.90268628663352...|
|  0.0|          0.0|[0.90268628663352...|
|  0.0|          0.0|[0.91813274435838...|
|  0.0|          0.0|[0.92035135496544...|
|  0.0|          0.0|[0.92035135496544...|

```

## Random Forest

```

# randomForest

rf_model = RandomForestClassifier()

# hyperparameter tuning
paramGrid = (
    ParamGridBuilder()
    .addGrid(rf_model.maxDepth, [5, 7, 10])
    .addGrid(rf_model.numTrees, [30,50,100])
    .build()
)

# Create a CrossValidator
evaluator = BinaryClassificationEvaluator()
cv = CrossValidator(
    estimator=rf_model,
    estimatorParamMaps=paramGrid,
    evaluator=evaluator,
    parallelism=8,
    numFolds=5,
)

cross_val_model = cv.fit(data_preprocessing_train)

# Extract the best model
best_model_rf = cross_val_model.bestModel
print(best_model_rf)

predictions_rf=best_model_rf.transform(data_preprocessing_test)

# evaluate best result
print( "RandomForestClassifier accuracy score on the test
set:",evaluator.evaluate(predictions_rf),"\\n")

print("RandomForestClassifier Prediction Table")
predictions_rf.select('label', 'prediction', 'probability').show(20)

RandomForestClassificationModel: uid=RandomForestClassifier_9dd4dd3636fa, num
Trees=100, numClasses=2, numFeatures=46
RandomForestClassifier accuracy score on the test set: 0.7786447919491395

RandomForestClassifier Prediction Table
+-----+-----+-----+
|label|prediction|          probability|

```

+	+	+	+
	0.0	0.0	[0.90150021148626...
	0.0	0.0	[0.90208502185240...
	0.0	0.0	[0.90152274208503...
	0.0	0.0	[0.92797679267685...
	0.0	0.0	[0.94355516985084...
	0.0	0.0	[0.92950653515829...
	0.0	0.0	[0.92475646105034...
	0.0	0.0	[0.94692517148005...
	0.0	0.0	[0.90534943231272...
	0.0	0.0	[0.90528069876450...
	0.0	0.0	[0.93199627291390...
	0.0	0.0	[0.93122995403098...

## Summary

We got close prediction results between Linear Regression, GBClassifier, and Random Forest model, however the GBClassifier is the best model for this project with the highest accuracy of 0.78.